

About the project

Case Study 6.1: NYC Taxi Trips

Case Study Description: To predict the trip duration of a New York taxi cab ride, we can build different types of features and evaluate them. We will start by describing what is a feature in this context; then we will develop some elementary features and add features using the software package `featuretools`. We will assess how these features perform in predicting trip duration.

Setup

Libraries

Note:

This uses an alternative to `featuretools` for R called `featuretoolsR`

- To Install `featuretoolsR` you need to have "devtools" installed.
- Then you can run: `devtools::install_github("magnusfurugard/featuretoolsR")`

In [1]:

```
# You will need to define (if not already) your Python path.  
# reticulate::py_discover_config()  
library("reticulate")  
use_python("~/Library/r-miniconda/envs/r-reticulate/bin/python")
```

In [2]:

```
library("featuretoolsR")  
library("dplyr")  
library("gbm")  
print("Functions loaded")
```

Registered S3 methods overwritten by 'ggplot2':

method	from
[.quosures	rlang
c.quosures	rlang
print.quosures	rlang

featuretoolsR 0.4.4

✓ Using Featuretools 0.16.0

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

Loaded gbm 2.1.5

```
[1] "Functions loaded"
```

Functions

- **featuretoolsR** does not include the original **encode_features** function, so a custom one was made called **custom_encode_features** which returns one-hot encoding of categorical variables.
- Additionally, custom function to replace python **SimpleImputer** class and **Rscores** were made.

In [3]:

```
# R^2 Calculation
r_squared <- function(predicted,actual) {
  r_squared <- 1 - (sum((actual - predicted)^2) / sum((actual - mean(actual))^2))
  return(r_squared)
}

# Custom encode feature to support missing function un featuretoolsR
custom_encode_features <- function(feature_matrix, to_encode, include_unknown, top_n) {
  temp <- feature_matrix
  for(feature in to_encode) {

    # We generate the list of top_n unique values for feature to encode
    if(include_unknown == TRUE) {
      encoded_feature <- names(head(sort(table(feature_matrix[[feature]],exclude = NULL),decreasing=TRUE),top_n))
    } else {
      encoded_feature <- names(head(sort(table(feature_matrix[[feature]]),decreasing=TRUE),top_n))
    }

    for(i in 1:length(encoded_feature)) {
      temp <- temp %>%
        mutate(
          # Create male column
          !!paste(feature,"=",encoded_feature[i]) := ifelse(feature_matrix[[feature]] == encoded_feature[i] , 1, 0)
        )
    }
  }
  return(temp[,!names(temp) %in% to_encode])
}

# Compute features function based on dsx-cs6 utils
compute_features <- function(features,entities) {

  # We create feature matrix using featuretools calculate_feature_matrix function
  feature_matrix <- calculate_feature_matrix(entities,features,
                                            approximate='36d',
                                            verbose=TRUE)

  # Since encode_features function is missing, we hot-encode pickup_neighborhood and
  dropoff_neighborhood
  # with a custom function created: custom_encode_features
  print("Finishing computing...")
  feature_matrix <- custom_encode_features(feature_matrix, to_encode=c("pickup_neighborhood","dropoff_neighborhood"), include_unknown=FALSE, top_n=1000)

  # Return Output
  return(feature_matrix)
}

# Logical to integer
logical_to_integer <- function(feature_matrix) {
  # Convert logical values to integers
  logical_features <- names(feature_matrix[,unlist(lapply(feature_matrix, is.logical))])
  for(feature in logical_features) {
    feature_matrix[[feature]] <- as.numeric(as.integer(feature_matrix[[feature]]))
  }
}
```

```

    return(feature_matrix)
  }

logical_to_factor <- function(feature_matrix) {
  # Convert logical values to integers
  logical_features <- names(feature_matrix[,unlist(lapply(feature_matrix, is.logical))])
  for(feature in logical_features) {
    feature_matrix[[feature]] <- as.factor(feature_matrix[[feature]])
  }
  return(feature_matrix)
}

custom_fit <- function(data) {
  out <- list()
  numeric_features <- names(data[,unlist(lapply(data, is.numeric))])
  for(feature in numeric_features) {
    out[[feature]] <- mean(data[[feature]] ,na.rm = TRUE)
  }
  return(out)
}

# Custom fit_transform
custom_transform <- function(data,fit) {
  numeric_features <- names(fit)
  for(feature in numeric_features) {
    data[[feature]][which(is.na(data[[feature]]))] <- fit[[feature]]
  }
  return(data)
}

# get_train_test_fm based on dsx-cs6 utils
get_train_test_fm <- function(feature_matrix, original_data, percentage) {

  out <- list()

  nrows <- nrow(feature_matrix)
  head <- nrows * percentage
  tail <- nrows - head

  X_train <- head(feature_matrix,head)
  y_train <- head(original_data$trip_duration,head)

  # Emulating fit_transform
  fit <- custom_fit(X_train)
  X_train = custom_transform(X_train,fit)

  X_test = tail(feature_matrix,tail)
  y_test = tail(original_data$trip_duration,tail)
  X_test = custom_transform(X_test,fit)

  # Return values
  out$X_train <- X_train
  out$y_train <- y_train
  out$X_test <- X_test
  out$y_test <- y_test

  return(out)
}
print("Functions Loaded")

```

```
[1] "Functions Loaded"
```

Data Load

In [4]:

```

# Needed to load pickle file
source_python("read_pickle.py")

```

- Opening a pickle file in R [additional info](#)

In [5]:

```
# We load data
trips <- read_pickle_file("./trips.pkl")
dropoff_neighborhoods <- read.csv("./dropoff_neighborhoods.csv")
pickup_neighborhoods <- read.csv("./pickup_neighborhoods.csv")

# We set type
trips$pickup_neighborhood <- as.character(trips$pickup_neighborhood)
trips$dropoff_neighborhood <- as.character(trips$dropoff_neighborhood)
dropoff_neighborhoods$neighborhood_id <- as.character(dropoff_neighborhoods$neighborhood_id)
pickup_neighborhoods$neighborhood_id <- as.character(pickup_neighborhoods$neighborhood_id)
print("Data Loaded")
```


[1] "Data Loaded"

In [6]:

```
# Histogram
hist(trips$trip_duration, main="Trip Duration Distribution",
     xlab="Trip Duration in Seconds",
     ylab="Number of Trips",
     col="lightblue")

# Tells us how many trips are in the dataset
paste("Trips in dataset:", nrow(trips))
```

'Trips in dataset: 10000'



Entities and Relationships

In [7]:

```
# Create entityset
entities <- as_entityset(
  trips,
  index = "id",
  entity_id = "trips",
  id = "trips",
  time_index = "pickup_datetime"
) %>%
add_entity(
  pickup_neighborhoods,
  entity_id = "pickup_neighborhoods",
```

```

    index = "neighborhood_id"
  ) %>%
  add_entity(
    dropoff_neighborhoods,
    entity_id = "dropoff_neighborhoods",
    index = "neighborhood_id"
  )

# Add index relationships
entities <- entities %>%
  add_relationship(
    parent_set = "pickup_neighborhoods",
    child_set = "trips",
    parent_idx = "neighborhood_id",
    child_idx = "pickup_neighborhood"
  ) %>%
  add_relationship(
    parent_set = "dropoff_neighborhoods",
    child_set = "trips",
    parent_idx = "neighborhood_id",
    child_idx = "dropoff_neighborhood"
  )
print("Entities and Relationships generated")

```

```
[1] "Entities and Relationships generated"
```

1.1 First Model

Transform Primitives

In [8]:

```

trans_primitives <- c("is_weekend")
features <- dfs(
  entityset=entities,
  target_entity="trips",
  trans_primitives=trans_primitives,
  agg_primitives=NULL,
  ignore_variables= list(trips=c("pickup_latitude", "pickup_longitude",
                                "dropoff_latitude", "dropoff_longitude", "trip_duration")),
  features_only=TRUE
)
print(paste("Number of features:",length(features)))

```

```
[1] "Number of features: 12"
```

Dataset Creation

In [9]:

```

feature_matrix <- compute_features(features, entities)
head(feature_matrix,5)

```

```
[1] "Finishing computing..."
```

A data.frame: 5 × 108

	vendor_id	passenger_count	trip_distance	payment_type	IS_WEEKEND(dropoff_datetime)	IS_WEEKEND(pickup_datetime)	pickup_ne
	<dbl>	<dbl>	<dbl>	<dbl>	<lgl>	<lgl>	
1	2	1	2.46	1	TRUE	TRUE	
2	1	2	7.90	1	TRUE	TRUE	
3	1	1	1.00	1	TRUE	TRUE	

4	2	1	0.02	2	TRUE	TRUE	
vendor_id	passenger_count	trip_distance	payment_type	IS_WEEKEND(dropoff_datetime)	IS_WEEKEND(pickup_datetime)	pickup_ne	
5	1	2	19.00	1	TRUE	TRUE	
edbs	edbs	edbs	edbs	edbs	edbs	edbs	

Split dataset

In [10]:

```
train_test = get_train_test_fm(feature_matrix, trips, .75)

y_train = log(train_test$y_train + 1)
y_test = log(train_test$y_test + 1)
X_train <- logical_to_integer(train_test$X_train)
X_test <- logical_to_integer(train_test$X_test)

trainingDataset = cbind(y_train,X_train)
print('Data split successful!')
```

[1] "Data split successful!"

Model 1.1: Fitting Generalized Boosted Regression Model

In [11]:

```
gbm <- gbm(y_train ~ ., data = trainingDataset,verbose = TRUE,
           n.trees = 100,
           distribution="gaussian",
           interaction.depth = 3,
           n.minobsinnode = 1,
           shrinkage = 0.1)
```

Warning message in gbm.fit(x = x, y = y, offset = offset, distribution = distribution, :
"variable 5: `IS_WEEKEND(dropoff_datetime)` has no variation."
Warning message in gbm.fit(x = x, y = y, offset = offset, distribution = distribution, :
"variable 6: `IS_WEEKEND(pickup_datetime)` has no variation."

Iter	TrainDeviance	ValidDeviance	StepSize	Improve
1	0.4814	nan	0.1000	0.0643
2	0.4293	nan	0.1000	0.0509
3	0.3847	nan	0.1000	0.0443
4	0.3480	nan	0.1000	0.0367
5	0.3179	nan	0.1000	0.0310
6	0.2921	nan	0.1000	0.0253
7	0.2710	nan	0.1000	0.0208
8	0.2526	nan	0.1000	0.0193
9	0.2379	nan	0.1000	0.0153
10	0.2255	nan	0.1000	0.0124
20	0.1603	nan	0.1000	0.0029
40	0.1348	nan	0.1000	0.0004
60	0.1278	nan	0.1000	0.0002
80	0.1243	nan	0.1000	0.0001
100	0.1211	nan	0.1000	0.0000

Model 1.1: Results from predictions

R^2

In [12]:

```
predicted_gbm <- predict(gbm,X_test,n.trees = 100)
paste("R Squared:",round(r_squared(y_test,predicted_gbm),4))
```

'R Squared': 0.6776'

R Squared: 0.0770

Model 1.1: Feature influence

In [13]:

```
head(round(relative.influence(gbm, sort= TRUE)/sum(relative.influence(gbm)), 4), 25)
```

n.trees not given. Using 100 trees.

n.trees not given. Using 100 trees.

```
trip_distance: 0.9102 dropoff_neighborhoods.longitude: 0.0336 dropoff_neighborhoods.latitude: 0.0213
pickup_neighborhoods.longitude: 0.0078 pickup_neighborhoods.latitude: 0.0049 `dropoff_neighborhood = AA`: 0.003
payment_type: 0.0022 `pickup_neighborhood = AR`: 0.0016 `dropoff_neighborhood = AU`: 0.0014
`pickup_neighborhood = AU`: 0.0012 `dropoff_neighborhood = AO`: 0.0012 vendor_id: 0.0011
`dropoff_neighborhood = AC`: 0.0011 `dropoff_neighborhood = H`: 9e-04 `pickup_neighborhood = D`: 9e-04
`pickup_neighborhood = AB`: 9e-04 `pickup_neighborhood = AG`: 8e-04 `pickup_neighborhood = X`: 7e-04
`dropoff_neighborhood = AR`: 6e-04 `pickup_neighborhood = AN`: 6e-04 `pickup_neighborhood = S`: 5e-04
passenger_count: 4e-04 `dropoff_neighborhood = AP`: 4e-04 `dropoff_neighborhood = J`: 4e-04
`dropoff_neighborhood = E`: 4e-04
```

1.2 Second Model

Transform Primitives

In [14]:

```
trans_primitives <- c("minute", "day", "week", "month", "weekday", "is_weekend")
features <- dfs(
  entityset=entities,
  target_entity="trips",
  trans_primitives=trans_primitives,
  agg_primitives=NULL,
  ignore_variables= list(trips=c("pickup_latitude", "pickup_longitude",
                                "dropoff_latitude", "dropoff_longitude", "trip_duration")),
  features_only=TRUE
)
print(paste("Number of features:", length(features)))
```

[1] "Number of features: 22"

Dataset Creation

In [15]:

```
feature_matrix <- compute_features(features, entities)
head(feature_matrix, 5)
```

[1] "Finishing computing..."

A data.frame: 5 × 118

	vendor_id	passenger_count	trip_distance	payment_type	MINUTE(dropoff_datetime)	MINUTE(pickup_datetime)	DAY(dropoff_datetime)
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	2	1	2.46	1	17	0	
2	1	2	7.90	1	24	0	
3	1	1	1.00	1	19	0	

4	2	1	0.02	2	1	0
vendor_id	passenger_count	trip_distance	payment_type	MINUTE(dropoff_datetime)	MINUTE(pickup_datetime)	DAY(dropoff_datetime)
5	1	2	19.00	1	58	1

Split dataset

In [16]:

```
train_test = get_train_test_fm(feature_matrix, trips, .75)

y_train = log(train_test$y_train + 1)
y_test = log(train_test$y_test + 1)
X_train <- logical_to_integer(train_test$X_train)
X_test <- logical_to_integer(train_test$X_test)

trainingDataset = cbind(y_train,X_train)
print('Data split successful!')
```

[1] "Data split successful!"

Model 1.2: Fitting Generalized Boosted Regression Model

In [17]:

```
gbm <- gbm(y_train ~ ., data = trainingDataset, verbose = TRUE,
           n.trees = 100,
           distribution="gaussian",
           interaction.depth = 3,
           n.minobsinnode = 1,
           shrinkage = 0.1)
```

Warning message in gbm.fit(x = x, y = y, offset = offset, distribution = distribution, :
"variable 9: `WEEK(dropoff_datetime)` has no variation."
Warning message in gbm.fit(x = x, y = y, offset = offset, distribution = distribution, :
"variable 10: `WEEK(pickup_datetime)` has no variation."
Warning message in gbm.fit(x = x, y = y, offset = offset, distribution = distribution, :
"variable 11: `MONTH(dropoff_datetime)` has no variation."
Warning message in gbm.fit(x = x, y = y, offset = offset, distribution = distribution, :
"variable 12: `MONTH(pickup_datetime)` has no variation."
Warning message in gbm.fit(x = x, y = y, offset = offset, distribution = distribution, :
"variable 15: `IS_WEEKEND(dropoff_datetime)` has no variation."
Warning message in gbm.fit(x = x, y = y, offset = offset, distribution = distribution, :
"variable 16: `IS_WEEKEND(pickup_datetime)` has no variation."

Iter	TrainDeviance	ValidDeviance	StepSize	Improve
1	0.4812	nan	0.1000	0.0629
2	0.4299	nan	0.1000	0.0522
3	0.3862	nan	0.1000	0.0441
4	0.3509	nan	0.1000	0.0358
5	0.3208	nan	0.1000	0.0315
6	0.2956	nan	0.1000	0.0246
7	0.2743	nan	0.1000	0.0218
8	0.2557	nan	0.1000	0.0186
9	0.2408	nan	0.1000	0.0150
10	0.2270	nan	0.1000	0.0137
20	0.1623	nan	0.1000	0.0039
40	0.1346	nan	0.1000	0.0005
60	0.1279	nan	0.1000	0.0003
80	0.1236	nan	0.1000	-0.0002
100	0.1210	nan	0.1000	0.0000

Model 1.2: Results from predictions

In [18]:

```
predicted_gbm <- predict(gbm,X_test,n.trees = 100)
paste("R Squared:",round(r_squared(y_test,predicted_gbm),4))
```

'R Squared: 0.7085'

Model 1.2: Feature influence

In [19]:

```
head(round(relative.influence(gbm,sort= TRUE)/sum(relative.influence(gbm)),4),25)
```

n.trees not given. Using 100 trees.

n.trees not given. Using 100 trees.

```
trip_distance: 0.9057 dropoff_neighborhoods.longitude: 0.0358 dropoff_neighborhoods.latitude: 0.02
pickup_neighborhoods.longitude: 0.0072 pickup_neighborhoods.latitude: 0.0048 payment_type: 0.0025
`dropoff_neighborhood = AA`: 0.0024 `DAY(pickup_datetime)`: 0.0022 `pickup_neighborhood = X`: 0.0019
`dropoff_neighborhood = AU`: 0.0018 `dropoff_neighborhood = AC`: 0.0014 `DAY(dropoff_datetime)`: 0.0012
`dropoff_neighborhood = E`: 0.0011 `pickup_neighborhood = AN`: 9e-04 `pickup_neighborhood = AR`: 9e-04
`MINUTE(dropoff_datetime)`: 7e-04 `pickup_neighborhood = AB`: 7e-04 `pickup_neighborhood = S`: 7e-04
`dropoff_neighborhood = D`: 6e-04 `dropoff_neighborhood = H`: 6e-04 `pickup_neighborhood = AG`: 5e-04
`pickup_neighborhood = AP`: 5e-04 `dropoff_neighborhood = AP`: 4e-04 `dropoff_neighborhood = J`: 4e-04
`MINUTE(pickup_datetime)`: 4e-04
```

1.3 Third Model

Transform Primitives

In [20]:

```
trans_primitives <- c("minute","day","week","month","weekday","is_weekend")
aggregation_primitives = c("count","sum","mean","median","std","max","min")
features <- dfs(
  entityset=entities,
  target_entity="trips",
  trans_primitives=trans_primitives,
  agg_primitives=aggregation_primitives,
  ignore_variables= list(trips=c("pickup_latitude", "pickup_longitude",
                                "dropoff_latitude", "dropoff_longitude", "trip_duration")),
  features_only=TRUE
)
print(paste("Number of features:",length(features)))
```

[1] "Number of features: 72"

Dataset Creation

In [21]:

```
feature_matrix <- compute_features(features, entities)
head(feature_matrix,5)
```

[1] "Finishing computing..."

A data.frame: 5 × 168

	vender_id	passenger_count	trip_distance	payment_type	MINUTE(dropoff_datetime)	MINUTE(pickup_datetime)	DAY(dropoff_datetime)
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	2	1	2.46	1	17	0	
2	1	2	7.90	1	24	0	
3	1	1	1.00	1	19	0	
4	2	1	0.02	2	1	0	
5	1	2	19.00	1	58	1	

Split dataset

In [22]:

```
train_test = get_train_test_fm(feature_matrix, trips, .75)

y_train = log(train_test$y_train + 1)
y_test = log(train_test$y_test + 1)
X_train <- logical_to_integer(train_test$X_train)
X_test <- logical_to_integer(train_test$X_test)

trainingDataset = cbind(y_train,X_train)
print('Data split successful!')
```

```
[1] "Data split successful!"
```

Model 1.3: Fitting Generalized Boosted Regression Model

In [23]:

```
gbm <- gbm(y_train ~ ., data = trainingDataset, verbose = TRUE,
           n.trees = 100,
           distribution="gaussian",
           interaction.depth = 3,
           n.minobsinnode = 1,
           shrinkage = 0.1)
```

```
Warning message in gbm.fit(x = x, y = y, offset = offset, distribution = distribution, :
"variable 9: `WEEK(dropoff_datetime)` has no variation."
Warning message in gbm.fit(x = x, y = y, offset = offset, distribution = distribution, :
"variable 10: `WEEK(pickup_datetime)` has no variation."
Warning message in gbm.fit(x = x, y = y, offset = offset, distribution = distribution, :
"variable 11: `MONTH(dropoff_datetime)` has no variation."
Warning message in gbm.fit(x = x, y = y, offset = offset, distribution = distribution, :
"variable 12: `MONTH(pickup_datetime)` has no variation."
Warning message in gbm.fit(x = x, y = y, offset = offset, distribution = distribution, :
"variable 15: `IS_WEEKEND(dropoff_datetime)` has no variation."
Warning message in gbm.fit(x = x, y = y, offset = offset, distribution = distribution, :
"variable 16: `IS_WEEKEND(pickup_datetime)` has no variation."
Warning message in gbm.fit(x = x, y = y, offset = offset, distribution = distribution, :
"variable 55: `dropoff_neighborhoods.MEDIAN(trips.passenger_count)` has no variation."
Warning message in gbm.fit(x = x, y = y, offset = offset, distribution = distribution, :
"variable 65: `dropoff_neighborhoods.MAX(trips.vendor_id)` has no variation."
Warning message in gbm.fit(x = x, y = y, offset = offset, distribution = distribution, :
"variable 67: `dropoff_neighborhoods.MIN(trips.passenger_count)` has no variation."
Warning message in gbm.fit(x = x, y = y, offset = offset, distribution = distribution, :
"variable 69: `dropoff_neighborhoods.MIN(trips.vendor_id)` has no variation."
Warning message in gbm.fit(x = x, y = y, offset = offset, distribution = distribution, :
"variable 70: `dropoff_neighborhoods.MIN(trips.payment_type)` has no variation."
```

Iter	TrainDeviance	ValidDeviance	StepSize	Improve
1	0.4813	nan	0.1000	0.0646
2	0.4281	nan	0.1000	0.0519
3	0.3847	nan	0.1000	0.0430
4	0.3482	nan	0.1000	0.0363
5	0.3164	nan	0.1000	0.0296

6	0.2912	nan	0.1000	0.0249
7	0.2704	nan	0.1000	0.0211
8	0.2529	nan	0.1000	0.0174
9	0.2370	nan	0.1000	0.0137
10	0.2249	nan	0.1000	0.0120
20	0.1621	nan	0.1000	0.0028
40	0.1345	nan	0.1000	0.0003
60	0.1266	nan	0.1000	0.0001
80	0.1214	nan	0.1000	0.0002
100	0.1187	nan	0.1000	-0.0001

Model 1.3: Results from predictions

R^2

In [24]:

```
predicted_gbm <- predict(gbm,X_test,n.trees = 100)
paste("R Squared:",round(r_squared(y_test,predicted_gbm),4))
```

'R Squared: 0.7137'

Model 1.3: Feature influence

In [25]:

```
head(round(relative.influence(gbm,sort= TRUE)/sum(relative.influence(gbm)),4),25)
```

n.trees not given. Using 100 trees.
n.trees not given. Using 100 trees.

```
trip_distance: 0.8968 dropoff_neighborhoods.longitude: 0.0274 dropoff_neighborhoods.latitude: 0.0181
`MINUTE(dropoff_datetime)`: 0.0083 `pickup_neighborhoods.MEAN(trips.passenger_count)`: 0.0045
`dropoff_neighborhoods.SUM(trips.trip_distance)`: 0.0045 `pickup_neighborhoods.COUNT(trips)`: 0.0042
`pickup_neighborhoods.SUM(trips.trip_distance)`: 0.0039 pickup_neighborhoods.longitude: 0.0029
`DAY(pickup_datetime)`: 0.0025 `dropoff_neighborhoods.MAX(trips.trip_distance)`: 0.0021
`dropoff_neighborhoods.SUM(trips.vendor_id)`: 0.0018 `pickup_neighborhood = AU`: 0.0016
`dropoff_neighborhoods.SUM(trips.passenger_count)`: 0.0015 `pickup_neighborhoods.MEAN(trips.trip_distance)`: 0.0014
payment_type: 0.0014 `dropoff_neighborhoods.MEDIAN(trips.trip_distance)`: 0.0014
`dropoff_neighborhoods.MEAN(trips.vendor_id)`: 0.0012 `dropoff_neighborhoods.MEAN(trips.trip_distance)`: 0.0012
pickup_neighborhoods.latitude: 0.0011 `dropoff_neighborhood = AA`: 9e-04
`dropoff_neighborhoods.SUM(trips.payment_type)`: 9e-04 `pickup_neighborhoods.MEDIAN(trips.payment_type)`: 8e-04
`pickup_neighborhoods.MEDIAN(trips.trip_distance)`: 8e-04 `pickup_neighborhoods.MAX(trips.trip_distance)`: 7e-04
```

Evaluate on Test Data

In [26]:

```
y_pred <- predict(gbm,X_test,n.trees = 100)
y_pred <- exp(y_pred) - 1 # undo the log we took earlier

y_test <- train_test$y_test

print('y_pred and y_test computation successful!')
```

[1] "y_pred and y_test computation successful!"

In [27]:

```
# Print the first 5 predictions and real values
```

```
head(round(y_pred),10)
head(round(y_test),10)
```

505 · 764 · 766 · 705 · 508 · 1531 · 520 · 1219 · 805 · 661

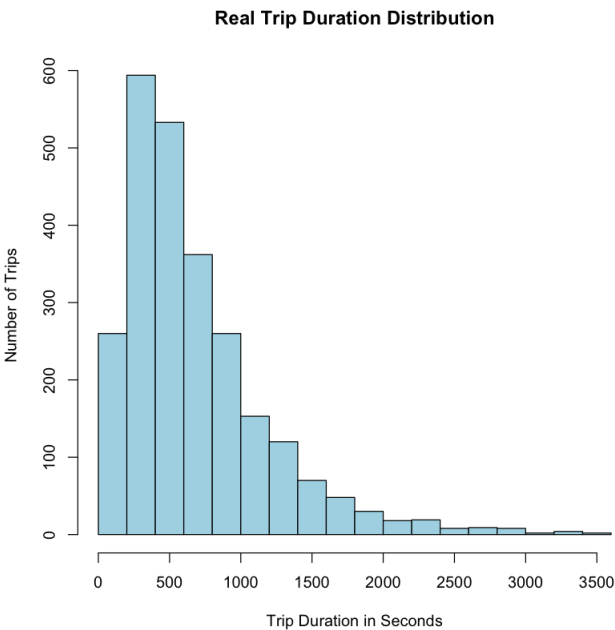
357 · 570 · 520 · 519 · 390 · 1146 · 553 · 1050 · 603 · 599

In [28]:

```
# Histogram of y_test
hist(y_test, main="Real Trip Duration Distribution",
     xlab="Trip Duration in Seconds",
     ylab="Number of Trips",
     col="lightblue")
summary(y_test)

# Histogram of y_pred
hist(y_pred, main="Predicted Trip Duration Distribution",
     xlab="Trip Duration in Seconds",
     ylab="Number of Trips",
     col="lightgreen")
summary(y_pred)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
2.0	321.0	546.5	681.0	885.2	3573.0



Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
12.63	356.67	583.79	700.14	956.31	2361.49

